

ML_models

January 3, 2021

1 ML model techniques

```
[1]: #importing necessary libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
[2]: dataframe=pd.read_csv(r"Financial/train_data.csv")
```

```
[3]: dataframe.head()
```

```
[3]:
```

	Time	V1	V2	V3	V4	V5	V6	\
0	38355.0	1.043949	0.318555	1.045810	2.805989	-0.561113	-0.367956	
1	22555.0	-1.665159	0.808440	1.805627	1.903416	-0.821627	0.934790	
2	2431.0	-0.324096	0.601836	0.865329	-2.138000	0.294663	-1.251553	
3	86773.0	-0.258270	1.217501	-0.585348	-0.875347	1.222481	-0.311027	
4	127202.0	2.142162	-0.494988	-1.936511	-0.818288	-0.025213	-1.027245	

	V7	V8	V9	...	V21	V22	V23	V24	\
0	0.032736	-0.042333	-0.322674	...	-0.240105	-0.680315	0.085328	0.684812	
1	-0.824802	0.975890	1.747469	...	-0.335332	-0.510994	0.035839	0.147565	
2	1.072114	-0.334896	1.071268	...	0.012220	0.352856	-0.341505	-0.145791	
3	1.073860	-0.161408	0.200665	...	-0.424626	-0.781158	0.019316	0.178614	
4	-0.151627	-0.305750	-0.869482	...	0.010115	0.021722	0.079463	-0.480899	

	V25	V26	V27	V28	Amount	Class
0	0.318620	-0.204963	0.001662	0.037894	49.67	0
1	-0.529358	-0.566950	-0.595998	-0.220086	16.94	0
2	0.094194	-0.804026	0.229428	-0.021623	1.00	0
3	-0.315616	0.096665	0.269740	-0.020635	10.78	0
4	0.023846	-0.279076	-0.030121	-0.043888	39.96	0

[5 rows x 31 columns]

2 Exploratory Data Analysis


```
[4]: #checking for standard deviation
dataframe.std()
```

```
[4]: Time      47500.410602
     V1         1.963028
     V2         1.661178
     V3         1.516107
     V4         1.415061
     V5         1.367074
     V6         1.325341
     V7         1.220384
     V8         1.192648
     V9         1.097367
     V10        1.087268
     V11        1.021904
     V12        0.999581
     V13        0.995449
     V14        0.959575
     V15        0.916011
     V16        0.875795
     V17        0.851222
     V18        0.838685
     V19        0.812614
     V20        0.772535
     V21        0.734187
     V22        0.724544
     V23        0.625165
     V24        0.606012
     V25        0.521348
     V26        0.482314
     V27        0.400286
     V28        0.331184
     Amount     248.100141
     Class       0.041548
     dtype: float64
```

```
[5]: #checking for correlation between Amount, Class and Time
df=dataframe.iloc[:,-2:]
df=df.join(dataframe["Time"])
sns.heatmap(df.corr(),annot=True,cmap="Blues")
print(df.corr())
```


	Amount	Class	Time
Amount	1.000000	0.006506	-0.011328
Class	0.006506	1.000000	-0.012800

Time -0.011328 -0.012800 1.000000



```
[6]: #checking for type of dataset Balanced or Imbalanced
X_imb=dataframe.iloc[:,0:-1]
Y_imb=dataframe.iloc[:,-1]
Y_imb.value_counts().plot.bar()
print(Y_imb.value_counts())
```

```
0    227451
1      394
Name: Class, dtype: int64
```



```
[7]: #Upscaling dataset
from imblearn.over_sampling import SMOTE
smk=SMOTE()
X,Y=smk.fit_sample(X_imb,Y_imb)
print(print(X_imb.shape,Y_imb.shape))
print(X.shape,Y.shape)
print(Y.value_counts())
Y.value_counts().plot.bar()
```

Using TensorFlow backend.

(227845, 30) (227845,)

None


(454902, 30) (454902,)

1 227451

0 227451

Name: Class, dtype: int64

[7]: <AxesSubplot:>



3 Modelling Techniques

```
[8]: #Importing Test Dataset
Test_data=pd.read_csv(r"Financial/test_data_hidden.csv")
X_train,Y_train=X,Y
X_test,Y_test=Test_data.iloc[:,0:-1],Test_data.iloc[:,-1]
```

```
[9]: pd.DataFrame(X_train)
```

```
[9]:
```

	Time	V1	V2	V3	V4	V5 \
0	38355.000000	1.043949	0.318555	1.045810	2.805989	-0.561113
1	22555.000000	-1.665159	0.808440	1.805627	1.903416	-0.821627
2	2431.000000	-0.324096	0.601836	0.865329	-2.138000	0.294663
3	86773.000000	-0.258270	1.217501	-0.585348	-0.875347	1.222481
4	127202.000000	2.142162	-0.494988	-1.936511	-0.818288	-0.025213
...
454897	93883.390713	-12.776379	7.324013	-17.364823	10.376465	-11.444020
454898	139850.607115	-1.043213	-0.617534	-2.982536	0.413148	0.919739
454899	158699.950195	-5.762649	-6.854235	-5.428864	5.146874	4.522205
454900	93885.053560	-12.149664	7.059878	-17.187042	9.271551	-10.772584
454901	47303.846632	-0.652141	1.198492	0.887867	1.223396	0.899199

	V6	V7	V8	V9	...	V20	V21 \
0	-0.367956	0.032736	-0.042333	-0.322674	...	-0.084556	-0.240105

1	0.934790	-0.824802	0.975890	1.747469	...	-0.373759	-0.335332
2	-1.251553	1.072114	-0.334896	1.071268	...	-0.039868	0.012220
3	-0.311027	1.073860	-0.161408	0.200665	...	0.382305	-0.424626
4	-1.027245	-0.151627	-0.305750	-0.869482	...	0.106592	0.010115
...
454897	-3.445067	-14.752594	8.045089	-5.562879	...	-1.166358	2.466291
454898	-0.901843	-0.376107	-0.350463	0.075369	...	-0.192415	1.393949
454899	-5.463446	-4.242289	0.864730	-1.155935	...	3.052142	1.431798
454900	-3.932420	-14.882749	7.138916	-5.061796	...	-0.875352	2.339961
454901	-0.361432	0.652013	0.093373	-0.979738	...	-0.137472	0.005726

	V22	V23	V24	V25	V26	V27	V28 \
0	-0.680315	0.085328	0.684812	0.318620	-0.204963	0.001662	0.037894
1	-0.510994	0.035839	0.147565	-0.529358	-0.566950	-0.595998	-0.220086
2	0.352856	-0.341505	-0.145791	0.094194	-0.804026	0.229428	-0.021623
3	-0.781158	0.019316	0.178614	-0.315616	0.096665	0.269740	-0.020635
4	0.021722	0.079463	-0.480899	0.023846	-0.279076	-0.030121	-0.043888
...
454897	-0.280985	-0.255514	0.544748	-0.682898	-0.210926	-2.234725	-0.627462
454898	1.829746	0.495892	0.054918	-0.518278	-0.104928	0.451686	0.083999
454899	-0.143725	0.814622	-0.278181	-1.038293	-0.679673	0.576835	-0.984923
454900	0.173664	-0.503640	0.537976	-1.032623	-0.243717	-1.815792	-0.436306
454901	-0.034321	-0.122503	-0.157211	-0.050876	-0.284419	-0.027712	0.072404

	Amount
0	49.670000
1	16.940000
2	1.000000
3	10.780000
4	39.960000
...	...
454897	31.648289
454898	185.907282
454899	286.496729
454900	20.452786
454901	3.766716

[454902 rows x 30 columns]

```
[10]: #creating a function to fit the models and display the results
def model_training(lists,X,Y,X_test,Y_test):
    from sklearn.linear_model import LogisticRegression
    from sklearn import svm
    from sklearn.naive_bayes import GaussianNB
    import xgboost
    from sklearn.ensemble import RandomForestClassifier
    from sklearn.metrics import accuracy_score,f1_score,roc_auc_score
```

```

Accuracy_score=[]
Auc_score=[]
F1_score=[]
models=[]
for k,v in lists.items():
    models.append(k)
    clf=v()
    #print(clf)
    clf.fit(X,Y)
    print("fitted",k)
    Y_pred=clf.predict(X_test)
    Accuracy_score.append(accuracy_score(Y_test,Y_pred))
    F1_score.append(f1_score(Y_test,Y_pred))
    Auc_score.append(roc_auc_score(Y_test,Y_pred))
    dic={'Models':models,"Accuracy":Accuracy_score,"f1_score":
    ↪F1_score,"AUC_score":Auc_score}
    return pd.DataFrame(dic)

```

```

[11]: from sklearn.linear_model import LogisticRegression
      from sklearn import svm
      from sklearn.naive_bayes import GaussianNB
      import xgboost
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import accuracy_score,f1_score,roc_auc_score

```

```

[49]: #standardizing the upscaled balanced training data sets
      from sklearn.preprocessing import StandardScaler
      sc=StandardScaler()
      X_train=sc.fit_transform(X_train)
      X_test=sc.transform(X_test)

```

```

[49]: array([[ 1.13050000e+05,  1.14697289e-01,  7.96303317e-01, ...,
               2.42724347e-01,  8.57125564e-02,  8.90000000e-01],
              [ 2.66670000e+04, -3.93177045e-02,  4.95783912e-01, ...,
               1.13135574e-01,  2.56835743e-01,  8.50000000e+01],
              [ 1.59519000e+05,  2.27570609e+00, -1.53150791e+00, ...,
               1.41970058e-02, -5.12892486e-02,  4.27000000e+01],
              ...,
              [ 1.38634000e+05,  2.20686736e+00, -7.48559259e-01, ...,
               -4.62001550e-02, -7.25858535e-02,  1.00000000e+00],
              [ 5.39070000e+04,  1.43057893e+00, -8.42353562e-01, ...,
               3.70949313e-02,  2.91797404e-02,  3.00000000e+01],
              [ 6.63730000e+04, -7.79271230e+00,  5.59993720e+00, ...,
               1.66156929e-01,  9.96801592e-02,  8.39000000e+00]])

```

```
[13]: #creating dictionary of the models and training the models using the balanced
      ↪data set

      #print(X_train.shape,Y_train.shape)
      #print(X_test.shape,Y_test.shape)
      model_list={'Logistic':LogisticRegression,"Naive Bayes":GaussianNB,'SVM':svm.
      ↪SVC,"Random Forest":RandomForestClassifier,'XGboost':xgboost.XGBClassifier}
      result=model_training(model_list,X_train,Y_train,X_test,Y_test)
      result
```

```
/usr/local/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:940:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

fitted Logistic

fitted Naive Bayes

fitted SVM

fitted Random Forest

fitted XGboost

```
[13]:
```

	Models	Accuracy	f1_score	AUC_score
0	Logistic	0.990643	0.241821	0.929101
1	Naive Bayes	0.976827	0.108108	0.896715
2	SVM	0.997683	0.538462	0.891881
3	Random Forest	0.999526	0.854054	0.902991
4	XGboost	0.999421	0.830769	0.913125

```
[18]: # training the models and displaying the results using imbalanced and
      ↪non-standardized dataset
      #Importing Test Dataset
      Test_data=pd.read_csv(r"Financial/test_data_hidden.csv")
      X_test_imb,Y_test_imb=Test_data.iloc[:,0:-1],Test_data.iloc[:,-1]
      imb_list={"Random Forest":RandomForestClassifier,'XGboost':xgboost.
      ↪XGBClassifier}
      result_imb=model_training(imb_list,X_imb,Y_imb,X_test_imb,Y_test_imb)
      result_imb
```

fitted Random Forest

fitted XGboost


```
[18]:
```

	Models	Accuracy	f1_score	AUC_score
0	Random Forest	0.999456	0.820809	0.862210
1	XGboost	0.999508	0.839080	0.872423

```
[19]: #comparing ensemble moedels with non-ensemble models with non-standardised
↪imbalanced dataset
model_list={'Logistic':LogisticRegression,"Naive Bayes":GaussianNB,'SVM':svm.
↪SVC,"Random Forest":RandomForestClassifier,'XGboost':xgboost.XGBClassifier}
result_imb=model_training(model_list,X_imb,Y_imb,X_test_imb,Y_test_imb)
result_imb
```

```
/usr/local/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:940:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

```
fitted Logistic
```

```
fitted Naive Bayes
```

```
fitted SVM
```

```
fitted Random Forest
```

```
fitted XGboost
```

```
[19]:
```

	Models	Accuracy	f1_score	AUC_score
0	Logistic	0.998964	0.677596	0.816124
1	Naive Bayes	0.992872	0.225191	0.797793
2	SVM	0.998280	0.000000	0.500000
3	Random Forest	0.999456	0.818713	0.857116
4	XGboost	0.999508	0.839080	0.872423